

Spark DataFrames Project Exercise

August 12, 2025

1 Spark DataFrames Project Exercise

Let's get some quick practice with your new Spark DataFrame skills, you will be asked some basic questions about some stock market data, in this case Walmart Stock from the years 2012-2017. This exercise will just ask a bunch of questions, unlike the future machine learning exercises, which will be a little looser and be in the form of "Consulting Projects", but more on that later!

For now, just answer the questions and complete the tasks below.

Use the `walmart_stock.csv` file to Answer and complete the tasks below!

Start a simple Spark Session

```
[2]: import findspark
findspark.init('/home/ian/spark-4.0.0-bin-hadoop3')
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('PROJECT').getOrCreate()
```

```
25/08/12 17:56:42 WARN SparkSession: Using an existing Spark session; only
runtime SQL configurations will take effect.
```

Load the Walmart Stock CSV File, have Spark infer the data types.

```
[5]: walmart_df = spark.read.csv('walmart_stock.csv', inferSchema=True, header=True)
```

What are the column names?

```
[6]: walmart_df.columns
```

```
[6]: ['Date', 'Open', 'High', 'Low', 'Close', 'Volume', 'Adj Close']
```

What does the Schema look like?

```
[7]: walmart_df.printSchema()
```

```
root
 |-- Date: date (nullable = true)
 |-- Open: double (nullable = true)
 |-- High: double (nullable = true)
 |-- Low: double (nullable = true)
 |-- Close: double (nullable = true)
 |-- Volume: integer (nullable = true)
```

```
|-- Adj Close: double (nullable = true)
```

Print out the first 5 columns.

```
[8]: walmart_df.head(5)
```

```
[8]: [Row(Date=datetime.date(2012, 1, 3), Open=59.970001, High=61.060001,
Low=59.869999, Close=60.330002, Volume=12668800, Adj Close=52.619234999999996),
Row(Date=datetime.date(2012, 1, 4), Open=60.209998999999996, High=60.349998,
Low=59.470001, Close=59.709998999999996, Volume=9593300, Adj Close=52.078475),
Row(Date=datetime.date(2012, 1, 5), Open=59.349998, High=59.619999,
Low=58.369999, Close=59.419998, Volume=12768200, Adj Close=51.825539),
Row(Date=datetime.date(2012, 1, 6), Open=59.419998, High=59.450001,
Low=58.869999, Close=59.0, Volume=8069400, Adj Close=51.45922),
Row(Date=datetime.date(2012, 1, 9), Open=59.029999, High=59.549999,
Low=58.919998, Close=59.18, Volume=6679300, Adj Close=51.616215000000004)]
```

Use describe() to learn about the DataFrame.

```
[11]: walmart_df.describe().show()
```

25/08/12 17:59:28 WARN SparkStringUtils: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

```
+-----+-----+-----+-----+-----+
+-----+-----+-----+
|summary|          Open|          High|          Low|
Close|          Volume|          Adj Close|
+-----+-----+-----+-----+
+-----+-----+-----+
| count|          1258|          1258|          1258|
1258|          1258|          1258|
| mean| 72.35785375357709|72.83938807631165|
71.9186009594594|72.38844998012726|8222093.481717011|67.23883848728146|
| stddev|
6.76809024470826|6.768186808159218|6.744075756255496|6.756859163732991|
4519780.8431556|6.722609449996857|
| min|56.389998999999996|          57.060001|          56.299999|
56.419998|          2094900|          50.363689|
| max|          90.800003|          90.970001|          89.25|
90.470001|          80898100|84.914216000000001|
+-----+-----+-----+-----+
+-----+-----+-----+
```

1.1 Bonus Question!

There are too many decimal places for mean and stddev in the describe() dataframe. Format the numbers to just show up to two decimal places. Pay careful attention to

the datatypes that `.describe()` returns, we didn't cover how to do this exact formatting, but we covered something very similar. [Check this link for a hint](#) If you get stuck on this, don't worry, just view the solutions.

```
[12]: stats_df = walmart_df.describe()
```

```
[24]: stats_df.printSchema()

## convert to numeric
from pyspark.sql.types import DoubleType
cols_to_change = stats_df.columns[1:]

for col in cols_to_change:
    stats_df = stats_df.withColumn(col, stats_df[col].cast(DoubleType()))
```

```
root
|-- summary: string (nullable = true)
|-- Open: string (nullable = true)
|-- High: string (nullable = true)
|-- Low: string (nullable = true)
|-- Close: string (nullable = true)
|-- Volume: string (nullable = true)
|-- Adj Close: string (nullable = true)
```

```
[26]: ## check types now
stats_df.printSchema()
```

```
root
|-- summary: string (nullable = true)
|-- Open: double (nullable = true)
|-- High: double (nullable = true)
|-- Low: double (nullable = true)
|-- Close: double (nullable = true)
|-- Volume: double (nullable = true)
|-- Adj Close: double (nullable = true)
```

```
[28]: from pyspark.sql.functions import format_number, mean

for col in cols_to_change:
    stats_df = stats_df.withColumn(col, format_number(col, 2))

stats_df.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
|summary|  Open|   High|    Low|   Close|      Volume|Adj Close|
+-----+-----+-----+-----+-----+-----+-----+
```

count	1,258.00	1,258.00	1,258.00	1,258.00	1,258.00	1,258.00
mean	72.36	72.84	71.92	72.39	8,222,093.48	67.24
stddev	6.77	6.77	6.74	6.76	4,519,780.84	6.72
min	56.39	57.06	56.30	56.42	2,094,900.00	50.36
max	90.80	90.97	89.25	90.47	80,898,100.00	84.91

Create a new dataframe with a column called HV Ratio that is the ratio of the High Price versus volume of stock traded for a day.

```
[32]: walmart_df = walmart_df.withColumn('HV Ratio', walmart_df['High']/
      ↪walmart_df['Volume'])
```

What day had the Peak High in Price?

```
[42]: max_high = walmart_df.agg({'High': 'max'}).collect()[0][0]
      walmart_df.filter(walmart_df['High'] == max_high).select('Date').show()
```

Date
2015-01-13

What is the mean of the Close column?

```
[44]: walmart_df.select(mean('Close')).show()
```

avg(Close)
72.38844998012726

What is the max and min of the Volume column?

```
[51]: from pyspark.sql.functions import min, max

      walmart_df.select(min('Volume'), max('Volume')).show()
```

min(Volume)	max(Volume)
2094900	80898100

How many days was the Close lower than 60 dollars?

```
[54]: walmart_df.filter(walmart_df['Close'] < 60).show()
```

Date	Open	High	Low
Close	Volume	Adj Close	HV Ratio
2012-01-04	60.209998999999996	60.349998	
59.470001	59.709998999999996	9593300	52.078475 6.290848613094555E-6
2012-01-05	59.349998	59.619999	58.369999
59.419998	12768200	51.825539	4.669412994783916E-6
2012-01-06	59.419998	59.450001	58.869999
59.0	8069400	51.45922	7.367338463826307E-6
2012-01-09	59.029999	59.549999	58.919998
59.18	6679300	51.616215000000004	8.915604778943901E-6
2012-01-10	59.43	59.709998999999996	
58.98	59.040001000000004	6907300	51.494109 8.644477436914568E-6
2012-01-11	59.060001	59.529999	59.040001000000004
59.400002	6365600	51.808098	9.351828421515645E-6
2012-01-12	59.790001000000004	60.0	59.400002
59.5	7236400	51.895315999999994	8.29141562102703E-6
2012-01-13			
59.18	59.610001000000004	59.009997999999996	59.540001000000004
7729300	51.930203999999996	7.712212102001476E-6	
2012-01-17	59.869999	60.110001000000004	59.52
59.849998	8500000	52.200581	7.071764823529412E-6
2012-02-22	59.580002	59.900002	58.369999
58.599998	28630200	51.110343	2.092196421960028E-6
2012-02-23	58.59		
58.900002	58.209998999999996	58.540001000000004	14880300
51.058014	3.958253664240641E-6		
2012-02-24	58.75	58.950001	
58.5	58.790001000000004	9925900	51.276061 5.939008150394423E-6
2012-02-27	58.700001		
58.779999	58.290001000000004	58.459998999999996	12258800
50.988237	4.794922749371879...		
2012-02-28	58.439999	59.099998	58.349998
58.93	10761900	51.398167	5.491595164422639E-6
2012-02-29	58.84	59.330002	58.720001
59.080002	11484400	51.528997	5.166138587997632E-6
2012-03-01	59.360001000000004	59.419998	58.639998999999996
58.82	16283900	51.302226	3.64900288014542E-6
2012-03-02	58.990002000000004	59.279999	
58.799999	59.009997999999996	9848100	51.467940999999996 6.019435119464668E-6
2012-03-05	58.959998999999996	59.59	58.75
59.400002	9651000	51.808098	6.174489690187545E-6
2012-03-06	59.040001000000004	59.220001	58.75
58.970001	9057100	51.433056	6.538516854180698E-6
2012-03-			

```
07|59.1100010000000004|59.8600010000000004|59.1100010000000004|59.8600010000000004|1
4916900|          52.209305|4.012898189302067...|
```

```
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
```

only showing top 20 rows

```
[55]: walmart_df.filter(walmart_df['Close'] < 60).count()
```

```
[55]: 81
```

What percentage of the time was the High greater than 80 dollars ?

In other words, (Number of Days High>80)/(Total Days in the dataset)

```
[57]: total_days = walmart_df.count()
total_days
```

```
[57]: 1258
```

```
[58]: days_more_than_80 = walmart_df.filter(walmart_df['High'] > 80).count()
days_more_than_80
```

```
[58]: 115
```

```
[61]: print(f' {(days_more_than_80/total_days)*100:.2f}%')
```

```
9.14%
```

What is the Pearson correlation between High and Volume?

[Hint](#)

```
[62]: from pyspark.sql.functions import corr

walmart_df.select(
    corr("High", "Volume")).show()
```

```
+-----+
| corr(High, Volume)|
+-----+
|-0.3384326061737161|
+-----+
```

```
[ ]: ## it seems that the more the volume traded, the lower the high price
```

What is the max High per year?

```
[71]: from pyspark.sql.functions import year, max
```

```

## generate date column
walmart_df = walmart_df.withColumn('year', year('Date'))

walmart_df.groupby('year').max('High').show()

```

```

+----+-----+
|year|max(High)|
+----+-----+
|2015|90.970001|
|2013|81.370003|
|2014|88.089996|
|2012|77.599998|
|2016|75.190002|
+----+-----+

```

What is the average Close for each Calendar Month?

In other words, across all the years, what is the average Close price for Jan, Feb, Mar, etc... Your result will have a value for each of these months.

```

[72]: from pyspark.sql.functions import month

## generate month column
walmart_df = walmart_df.withColumn('month', month('Date'))

walmart_df.groupby('month').mean('Close').show()

```

```

+----+-----+
|month|      avg(Close)|
+----+-----+
| 12|72.84792478301885|
|  1|71.44801958415842|
|  6| 72.4953774245283|
|  3|71.77794377570092|
|  5|72.30971688679247|
|  9|72.18411785294116|
|  4|72.97361900952382|
|  8|73.02981855454546|
|  7|74.43971943925233|
| 10|71.57854545454543|
| 11| 72.1110893069307|
|  2| 71.306804443299|
+----+-----+

```

2 Great Job!

2.1 OTHER INTERESTING INSIGHTS

2.1.1 1. 7 day moving average

```
[73]: from pyspark.sql.window import Window
      from pyspark.sql.functions import avg, col
```

```
[82]: windowSpec = Window.orderBy('Date').rowsBetween(-6,0)
```

```
[83]: walmart_with_avg = walmart_df.withColumn(
      "7_day_moving_avg", avg(col("Close")).over(windowSpec)
    )
```

```
[88]: walmart_with_avg.select("Date", "Close", "7_day_moving_avg").show()
```

25/08/12 18:34:56 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance
degradation.

25/08/12 18:34:56 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance
degradation.

25/08/12 18:34:56 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance
degradation.

25/08/12 18:34:56 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance
degradation.

25/08/12 18:34:56 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance
degradation.

```
+-----+-----+-----+
|      Date|      Close| 7_day_moving_avg|
+-----+-----+-----+
|2012-01-03|      60.330002|      60.330002|
|2012-01-04|59.709998999999996|60.020000499999995|
|2012-01-05|      59.419998| 59.819999666666666|
|2012-01-06|      59.0|59.614999749999996|
|2012-01-09|      59.18|59.527999799999996|
|2012-01-10|59.040001000000004|59.446666666666665|
|2012-01-11|      59.400002| 59.44000028571429|
|2012-01-12|      59.5| 59.32142857142857|
|2012-01-13|59.540001000000004| 59.29714314285714|
|2012-01-17|      59.849998| 59.35857171428571|
|2012-01-18|60.009997999999996| 59.50285714285714|
|2012-01-19|60.610001000000004| 59.707143000000001|
|2012-01-20|61.009997999999996| 59.98857114285714|
```


2012-01-23	60.91	60.20428514285714
2012-01-24	61.389998999999996	60.474284999999995
2012-01-25	61.470001	60.74999928571429
2012-01-26	60.970001	60.909999714285725
2012-01-27	60.709998999999996	61.00999985714286
2012-01-30	61.299999	61.108571
2012-01-31	61.360001000000004	61.15857142857143

+-----+-----+-----+

only showing top 20 rows

2.1.2 2. Year over year growth

```
[89]: from pyspark.sql.functions import year, lag, col
      from pyspark.sql.window import Window

      # Add a 'year' column to the DataFrame
      walmart_with_year = walmart_df.withColumn('year', year('Date'))

      # Get the last closing price of each year
      last_close_of_year = walmart_with_year.groupBy('year').agg({'Close': 'last'}).
      ↪withColumnRenamed('last(Close)', 'Close')

      # Use a window function to get the previous year's close price
      windowSpec_year = Window.orderBy('year')
      last_close_of_year_with_prev = last_close_of_year.withColumn('prev_year_close',
      ↪lag(col('Close'), 1).over(windowSpec_year))

      # Calculate the year-over-year growth percentage
      yoy_growth = last_close_of_year_with_prev.withColumn(
          'yoy_growth_percent',
          ((col('Close') - col('prev_year_close')) / col('prev_year_close')) * 100
      )

      # Show the results
      yoy_growth.select('year', 'Close', 'yoy_growth_percent').show()
```

25/08/12 18:37:55 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance degradation.

25/08/12 18:37:55 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance degradation.

25/08/12 18:37:55 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance degradation.

25/08/12 18:37:55 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance degradation.

25/08/12 18:37:55 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance
degradation.

```
+-----+-----+-----+
|year|    Close|yoy_growth_percent|
+-----+-----+-----+
|2012|68.230003|          NULL|
|2013|78.690002|15.330497640458862|
|2014|85.879997| 9.137113759382032|
|2015|61.299999|-28.62133076227285|
|2016|69.120003|12.756939849215982|
+-----+-----+-----+
```

25/08/12 18:37:55 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance
degradation.

25/08/12 18:37:55 WARN WindowExec: No Partition Defined for Window operation!
Moving all data to a single partition, this can cause serious performance
degradation.

[]: